

# CSMT-GNN: Structure-Aware Contextual Variable Dropout for Code Language Models

Shengyao Liang

Independent Researcher

pikeshuaiwe@gmail.com

ORCID: <https://orcid.org/0009-0002-3713-8700>

June 2026

## Abstract

Code language models can suffer from *structural hallucination*: they produce syntactically fluent programs while silently breaking program dependencies, such as shadowed variables, invalid imports, or guards that no longer protect later accesses. To mitigate this failure mode, I propose CSMT-GNN, a structure-aware code language model prototype that keeps token generation primary while injecting tokenizer-aligned AST ids and prefix-masked block graph states. The architecture combines block-local Transformer computation, residual AST-gated pooling, autoregressive block-level graph message passing, graph-to-token injection, and Contextual Variable Dropout (CVD), a structured regularizer over definition-bearing blocks. I also treat the development process as part of the artifact: the first prototype was produced through multi-LLM assisted coding and then rewritten after audit, which exposed both the speed and the danger of that workflow. This release provides a runnable reference implementation, theoretical checks for information flow and training dynamics, an explicit cost boundary for the AST side channel, local mechanism diagnostics, and a staged evaluation plan. I do not claim frontier-scale results. I claim that the idea and the workflow are now honest enough to be tested. **Code and reproducible scripts:** <https://github.com/ShengyaoLiang/csmt-gnn>; **archived release:** <https://doi.org/10.5281/zenodo.20844185>.

## 1 Introduction

The failure mode that pushed me toward this project was not that code models cannot write syntax. They often can. The more unsettling failure is that a model can keep the syntax fluent while quietly losing the program relation that made the next line true: a variable was shadowed, a method belonged to a different object, an import was never in scope, or a guard no longer protected a later access. In my first notes I called this *causal hallucination*, because the visible mistake usually comes from a broken dependency. I now call it *structural hallucination* unless I am making an actual causal claim. That distinction matters. The generated code looks ordinary until the dependency that should have anchored it is gone.

My starting hypothesis is simple. If the task is code, then the model should not be forced to rediscover every bit of syntactic structure from raw tokens at every layer. The token stream should remain primary, because generation is still token generation. But an AST can tell the model, cheaply and repeatedly, what kind of syntactic object a token participates in. A block-level graph can then move summaries across the sequence at a granularity closer to statements, definitions, and uses. CSMT-GNN is the architecture I built around that hypothesis, and the rest of the draft is me trying to keep that hypothesis from outrunning the evidence.

There is also a correction of my own earlier framing. I first called the masking operation a causal do-intervention. That was too strong. In Pearl’s structural causal model, an intervention cuts the incoming mechanisms of a variable and replaces it with an externally set value. My code does not identify an SCM over program semantics, and it does not perform adjustment. It replaces selected value messages in a neural block graph. Writing that down was uncomfortable, but it made the work cleaner. The honest name is *Contextual*

*Variable Dropout* (CVD): a structured regularizer aimed at definition-bearing blocks. I would rather publish a smaller claim that can be tested than a larger one that collapses under inspection.

The public artifact has five concrete pieces. First, it defines a code-modeling architecture that combines block-local attention, trainable AST embeddings, prefix-masked block graph attention, gated graph-to-token injection, and top- $k$  mixture-of-experts (MoE) computation. Second, it ships a cleaned implementation that removes several failure modes in the earlier prototype: random non-trainable AST embeddings, token-offset mismatch, block-size-specific softmax code, distorted CVD sampling, and a Python-loop MoE path that collapses throughput when specialized kernels are unavailable. Third, it gives mathematical checks for prefix information flow, attention-edge count, AST gate stability, the cost boundary of the syntax channel, and the stochastic objective optimized by CVD. Fourth, it documents the multi-LLM assisted development workflow that produced the first draft and the audit steps that made the artifact defensible. Fifth, it gives the evaluation protocol I would trust before making a stronger claim.

Beyond the architecture, this project argues for a workflow and a stance. I care about code models because code is becoming a production material for science, engineering, data analysis, and automation. I am less interested in making AI look human than in making it a reliable part of human productivity. If generated code is fluent but structurally wrong, it does not merely fail as text; it damages a tool chain. That is why I designed CSMT-GNN around program structure rather than around a broader claim about reasoning in general.

I used multiple LLMs during early prototype development because I wanted to see whether an independent researcher could move from architecture thinking to a runnable artifact without institutional-scale support. Different systems helped with different parts of the work: scaffolding, implementation, criticism, and documentation. That made research feel less locked behind infrastructure: one person could explore preprocessors, kernels, training code, and documentation much faster.

But the same workflow produced mistakes that serious science cannot keep: overconfident causal language, repeated docs, brittle hard-coded kernels, random offline AST embeddings, and code names that sounded more rigorous than the implementation. My later work was audit and authorship: narrowing the claims, rewriting the code, removing the pseudo-causal frame, and making the artifact inspectable. For the arXiv and public GitHub versions, I keep the fuller author story because the methodological claim matters too. The lesson is not that AI tools make research automatic. It is that more people can start from a serious question, use AI to reach a testable artifact, and then take responsibility for the result.

I keep the claim deliberately narrow. CSMT-GNN is a structure-aware code-model architecture and a runnable research artifact; it is not a proof that AST features are always useful, not a Pearl-style causal model, and not a frontier-scale empirical result. CVD is treated as structured stochastic regularization rather than a formal do-intervention. The purpose of this release is to make the architecture, its assumptions, and its failure modes inspectable enough for others to test.

**Code and reproducible scripts:** <https://github.com/ShengyaoLiang/csmt-gnn>. The current v0.1.1 archived release is <https://doi.org/10.5281/zenodo.20844185>; the stable all-version Zenodo concept DOI is <https://doi.org/10.5281/zenodo.20840624>.

## 2 Related Work

**Language models for code.** Large Transformer models trained on code have become practical programming assistants [Chen et al., 2021, Nijkamp et al., 2023, Roziere et al., 2023]. Their success comes from scale, data, and the generality of next-token prediction. However, benchmarks such as HumanEval and MBPP largely measure functional success at the problem level and do not isolate structural failure modes. Work on execution feedback, retrieval, and program repair improves reliability, but many deployments still need the base model to preserve code state before a tool is invoked.

**Syntax-aware neural code models.** ASTs and program graphs have long been used for code representation learning, including tree-based neural networks, graph neural networks, and path-based representations [Allamanis et al., 2018, Alon et al., 2019, Brockschmidt et al., 2019]. CSMT-GNN does not replace token modeling with a graph model. Instead, it keeps a Transformer language model as the main computation and injects compact AST features into block pooling and graph communication.




Figure 1: One CSMT-GNN layer. Solid arrows are differentiable model paths; dashed arrows are structural side inputs or training-time regularization. CVD is applied only during training.

**Efficient Transformers and MoE.** The architecture borrows from efficient attention and expert routing ideas: local attention reduces the quadratic cost within a bounded block, compressed key-value projections reduce activation size, and MoE layers increase capacity without activating every parameter for every token [Vaswani et al., 2017, Dao et al., 2022, Shazeer et al., 2017, Lepikhin et al., 2021, Fedus et al., 2022]. The present code uses PyTorch scaled-dot-product attention for correctness and portability, while leaving room for fused kernels after the algorithmic contract is stable.

### 3 Method

Figure 1 is the most compact view of the layer. The token stream remains the main computation; AST features only modulate pooling, and the block graph communicates summaries rather than replacing token generation.

#### 3.1 Inputs, Offline AST Features, and Prefix Inference Features

Let a token sequence be  $x_{1:L}$  and let  $B$  be the block size. The sequence is padded to  $M = \lceil L/B \rceil$  blocks. The offline preprocessor parses each source file and writes three compact arrays:

$$A \in \mathbb{N}^{M \times B}, \quad r \in \{0, 1\}^M, \quad u \in \{0, 1\}^L,$$

where  $A_{m,b}$  is an AST node-type id for token  $b$  in block  $m$ , and  $r_m$  marks whether block  $m$  contains a variable-definition event. The token mask  $u_t$  marks definition tokens before they are reduced to block eligibility for CVD. During training, the model owns a shared embedding table  $E_{\text{ast}}$  and computes  $e_{m,b} = E_{\text{ast}}(A_{m,b})$ . This is a deliberate correction of the earlier prototype, which materialized random AST embeddings offline and therefore made the syntax channel non-trainable.

When a HuggingFace fast tokenizer is provided, the preprocessor uses its offset mapping so AST spans follow the same tokenization as the language-model tokens. Without that tokenizer, the preprocessor falls back to Python lexical tokens and marks the artifact accordingly; I treat that mode as a debugging path, not as a claim of BPE-level alignment. For offline training artifacts, the reference preprocessor now defaults to a complete-file parse, because this gives the cleanest AST type ids and the most reliable definition masks. The token-level definition mask is built from Python binding spans when the built-in parser can parse the file; tree-sitter remains the source of AST type ids and a best-effort fallback for incomplete code. Prefix and per-token prefix parsing remain explicit modes for leakage studies and generation-time feature construction, not the default training path.

Autoregressive inference cannot assume a complete, syntactically valid file. The reference code therefore includes a prefix feature builder that reparses the current generated prefix, relies on tree-sitter recovery

when possible, and falls back to lexical ids when the prefix is too broken to yield a useful AST node. At inference time, a saved AST vocabulary is frozen: unseen node types map to <UNKNOWN> rather than creating embedding ids that were never trained. The script reports the unknown-token rate alongside latency. This does not make inference free. It turns AST availability into an explicit systems question: measure parser latency, recovery quality, vocabulary fallback, and the degradation from offline features to prefix features. I do not claim that this cost is solved by the architecture.

### 3.2 Block-Local Transformer Computation

Token embeddings and positional embeddings form  $H^0 \in \mathbb{R}^{M \times B \times d}$ . Each layer first applies RMSNorm and block-local self-attention:

$$\tilde{H}_m = H_m + \text{Attn}_B(\text{RMSNorm}(H_m)).$$

The attention uses compressed key-value projections:

$$Q = HW_Q, \quad C = HW_C, \quad K = CW_K, \quad V = CW_V,$$

which keeps the interface close to multi-head latent attention while avoiding a custom kernel dependency in the reference implementation.

### 3.3 AST-Gated Block Pooling

The AST side channel modulates token states before pooling. I use a residual gate rather than a plain sigmoid gate so the initial model does not halve the hidden-state scale. With  $g_{m,b} = 1 + \lambda \tanh(W_{\text{ast}} e_{m,b})$ , the gated state is

$$\bar{H}_{m,b} = \tilde{H}_{m,b} \odot g_{m,b}.$$

Before pooling, an initially-zero boundary module lets the first  $w_{\text{bdry}}$  tokens of block  $m$  receive learned corrections from the last token of block  $m - 1$ . I keep this correction one-way so the language model does not read future tokens across a block boundary. This is a small concession to the fact that fixed blocks can cut through syntax. Because the boundary projection is initialized at zero, the initial model is exactly the no-boundary model; training must learn whether the correction is useful. A learned query vector  $q$  scores tokens within each block:

$$\alpha_{m,b} = \text{softmax}_b \left( \frac{q^\top \bar{H}_{m,b}}{\sqrt{d}} \right), \quad z_m = \sum_b \alpha_{m,b} \bar{H}_{m,b}.$$

The result  $z_m$  is a compact block state used by graph attention. This design keeps AST computation proportional to  $Ld_{\text{ast}}d$  in the unfused reference implementation and makes the cost visible for ablation.

### 3.4 Prefix Block Graph and Contextual Variable Dropout

The block states  $z_{1:M}$  are passed through a prefix-masked graph attention layer. It is graph-like because nodes are blocks and edges point from earlier blocks to later blocks; it is autoregressive because node  $m$  cannot attend to  $n > m$ :

$$s_{1:M} = \text{PrefixGraphAttn}(z_{1:M}).$$

In batched training, this prefix mask is combined with a valid-block key mask, and padded block outputs are zeroed after graph attention. I include this detail because padding should be outside the block graph, not merely harmless because later token states are masked.

CVD changes only the value messages. During training, each variable-definition block is sampled with probability  $p_{\text{cvd}}$ . The current implementation prefers token-level definition masks when available and reduces them to block eligibility; this reduces false positives from coarse block-level labels. For selected blocks, the value vector is replaced by a learned vector  $v_{\text{cvd}}$  before graph attention:

$$V_m = \begin{cases} v_{\text{cvd}}, & r_m = 1 \text{ and } u_m < p_{\text{cvd}}, \\ W_V z_m, & \text{otherwise.} \end{cases}$$

I no longer call this a causal intervention. The operation is a structured dropout mechanism targeted at definition-bearing blocks. The intended effect is to discourage brittle reliance on one local definition token and to encourage redundant, contextually grounded representations.

### 3.5 Graph-to-Token Injection and MoE

The graph state is injected into token states with a prefix shift, so block  $m$  receives graph context from previous blocks:

$$c_m = \begin{cases} 0, & m = 1, \\ W_c \text{ RMSNorm}(s_{m-1}), & m > 1. \end{cases}$$

The token-level injection is gated:

$$H'_{m,b} = \bar{H}_{m,b} + W_o (\sigma(W_g \bar{H}_{m,b}) \odot c_m).$$

Finally, the layer applies a residual top- $k$  MoE feed-forward module. The reference code sorts token-expert assignments into contiguous expert batches, which avoids the earlier per-token dynamic dispatch pattern and is a reasonable fallback when vendor grouped GEMM kernels are not available. The implementation also tracks a router load-balancing loss and a router z-loss; for production training, a grouped-GEMM expert kernel remains the right target.

## 4 Implementation and Development Workflow

### 4.1 Reference Code

My cleaned implementation is organized around small, auditable entry points:

- `ast_preprocessor.py` emits AST ids, block masks, token-level definition masks, and an AST vocabulary.
- `inference_ast.py` builds approximate AST ids from incomplete generation prefixes.
- `diagnostics.py` creates small shadowing, long-range import, and guarded-attribute cases for cheap pipeline checks.
- The structural-probe script measures whether those diagnostic cases actually contain definition-use and cross-block structure.
- The CVD audit script records how often variable-targeted and random CVD sample eligible blocks. This accounting is opt-in because it synchronizes device tensors for human-readable counts; the default training path samples masks without asking Python whether any block was selected.
- `transformer_baseline.py` implements a separate token-only causal Transformer, including a rough parameter-neighbor control used by the tiny diagnostics. It shares the same fail-fast token-id and length checks as the CSMT path, so baseline runs do not silently repair malformed inputs.
- `csmt_gnn.py` contains the reusable architecture, and `train.py` supports single-process and `torchrun` training over paired token and AST NumPy files. Input id range scans are enabled by default for fail-fast debugging, but can be disabled after the NumPy pipeline is validated; dtype, shape, and length checks remain active. The trainer also trims AST ids and definition masks to the next-token model-input prefix before each forward pass, so complete-sample side inputs are not handed to a shorter prediction prefix.
- `scripts/validate_data_pipeline.py` is the NumPy preflight check for that validated-pipeline fast path: it scans token ids, AST ids, block shapes, masks, and usable lengths before training.
- `scripts/diagnostic_poc_train.py` trains an ablation matrix on the diagnostic cases when PyTorch is available. It now trims AST ids and definition masks to the same next-token prefix before each CSMT forward pass and records a `prefix_feature_audit` field in its JSON output.

I treat the diagnostic output as a plumbing and falsification check, not as a benchmark.

## 4.2 Ablation-Ready Components

The reference model exposes each contested mechanism as a real configuration switch rather than a rhetorical promise. `use_ast_gate` removes the AST embedding table and AST projection from the model. `use_block_graph` removes the prefix graph and graph-to-token injection. `use_cvd` removes the CVD replacement parameter and sampling path. `use_boundary` removes the causal boundary mixer. `use_moe` replaces the routed MoE with a dense SwiGLU-style feed-forward module, so the no-MoE ablation still has feed-forward capacity. The CVD sampler also supports `cvd_scope=variable` and `cvd_scope=random`, which makes the ordinary-noise control explicit. These switches matter because I cannot ask a reader to trust an architecture whose main objections cannot be isolated. I also keep the Transformer baselines outside this switch system. The `transformer_baseline` variant is a separate causal Transformer with no AST, block graph, boundary mixer, or CVD path. The `transformer_matched` variant widens the Transformer feed-forward layer to be a rough parameter neighbor of the tiny CSMT variants. It is not a perfect matched study, but it is a stronger first control than comparing CSMT only against its own disabled configurations.

## 4.3 Development Audit

The author workflow described in the introduction leaves a concrete audit trail. The early prototype was useful because it made the architecture runnable, but it also exposed failure modes that would have made the work indefensible if kept: causal terminology that the computation did not earn, repeated operation guides, a block-size-specific softmax path, random offline AST embeddings, distorted CVD sampling, and a Python-loop expert path that would collapse throughput when specialized kernels were unavailable.

The current reference implementation is the result of removing those weaknesses rather than polishing their names. I renamed the unsupported do-intervention framing to CVD, moved AST handling to tokenizer-aligned ids with trainable embeddings, removed the brittle fused-kernel path from the reference implementation, added prefix inference features, corrected batch and padding semantics, added Transformer controls, and made the release package inspectable. This is why I treat the development process as part of the contribution. I am not presenting multi-LLM coding as a substitute for authorship; I am presenting it as a way for an independent researcher to reach a testable artifact faster, provided the final work is audited, renamed where necessary, and stripped of claims the implementation cannot support. The artifact is not only a proposed architecture, but also a record of which claims survived audit and which ones had to be deleted.

My choice to remove the hand-written Triton F1 kernel from the reference path is intentional. The earlier kernel assumed  $B = 64$  in its second-stage softmax. A production kernel should either be generated for supported block sizes or use a shape-parametric reduction. Until that exists, the correct reference implementation is more valuable than a fast path with undefined behavior at other block sizes.

## 5 Theoretical Checks and Complexity

Algebra cannot replace training evidence, but it can rule out basic failure modes before scaling: future leakage in the block graph, unbounded changes from the AST gate at initialization, unsupported memory claims, and a CVD story with no objective behind it. These checks are deliberately modest; their purpose is to make the mechanism auditable before any performance claim is attached to it. Detailed proofs are collected in Appendix A, so the main text can stay focused on what each statement means for the architecture.

**Proposition 1** (Prefix visibility). *Assume each block-local attention operator uses a causal mask, the block graph attends only from block  $m$  to valid blocks  $n \leq m$ , and padded block outputs are not injected into real token states. If the AST ids for token  $t$  are computed from a source prefix ending no later than  $t$ , then the hidden state used to predict token  $t + 1$  is a function only of tokens and AST spans in the prefix  $x_{1:t}$ .*

The default offline training preprocessor uses complete-file parsing for cleaner syntax labels and definition masks. Proposition 1 is therefore not a claim that every training artifact is prefix-only. It is the contract for generation-time features and for experiments that require strict prefix preprocessing. When training uses complete-file features and inference uses prefix features, the relevant question is the measured drift between those two side inputs, which I define in Proposition 7.

Table 1: Problems in the original prototype and their corresponding fixes.

| Area                  | Original risk                                         | Current fix                                                            |
|-----------------------|-------------------------------------------------------|------------------------------------------------------------------------|
| AST features          | Offline random embeddings                             | Offline ids, trainable shared embedding                                |
| Token offsets         | Lexical/BPE mismatch risk                             | Optional tokenizer offset mapping                                      |
| Causal claim          | Mask framed as do-calculus                            | Renamed and scoped as CVD regularization                               |
| Fused softmax         | Hard-coded four chunks of 16                          | Reference path uses dynamic block shape                                |
| CVD sampling          | Forced one mask per eligible sample                   | Bernoulli sampling with true zero-mask cases                           |
| CVD auditing          | Per-step count sync in training path                  | Audit counts are opt-in; default sampling is synchronization-free      |
| Padding in graph path | Padded blocks could remain graph nodes                | Valid-block attention mask and zeroed padded block outputs             |
| Trainer side inputs   | Full-sample AST features during next-token prediction | Prefix-trimmed AST ids and masks before forward pass                   |
| Inference AST         | Assumed complete-file parse                           | Prefix parser with lexical fallback                                    |
| Block boundaries      | Hard cuts through syntax                              | Initially-zero causal boundary mixing before pooling                   |
| MoE fallback          | Token-wise dynamic dispatch                           | Sorted top- $k$ contiguous expert batches                              |
| Ablations             | Only baseline/full comparison                         | Component switches plus random-dropout control                         |
| Documentation         | Repeated guides inside code                           | Separate source files and paper artifact                               |
| Training checks       | Per-step id range scans always on                     | Default fail-fast checks with an explicit validated-pipeline fast path |

**Proposition 2** (Causal boundary mixing). *Let the boundary module update only the first  $w_{\text{bdry}}$  token states of block  $m$  using a learned function of the last token state of block  $m - 1$ :*

$$\hat{H}_{m,b} = \tilde{H}_{m,b} + \mathbf{1}[b \leq w_{\text{bdry}}] \gamma P_b(\tilde{H}_{m-1,B}), \quad m > 1.$$

*If  $\tilde{H}_{m-1,B}$  is causal with respect to tokens up to the end of block  $m - 1$ , then  $\hat{H}_{m,b}$  is causal with respect to tokens up to position  $(m - 1)B + b$ . In particular, the boundary module cannot transmit information from block  $m$  or any later block into block  $m - 1$ .*

**Proposition 3** (Attention-edge count). *For  $L = MB$  tokens, full causal token attention uses  $L(L + 1)/2$  token-token edges. CSMT-GNN uses  $MB(B + 1)/2$  block-local token edges plus  $M(M + 1)/2$  block graph edges. The edge ratio is*

$$\rho(B, M) = \frac{MB(B + 1) + M(M + 1)}{MB(MB + 1)}.$$

*For fixed  $B$  and large  $M$ , the graph term contributes an asymptotic ratio  $1/B^2$  and the local term contributes  $O(1/M)$ . Thus a fixed block size gives a constant-factor reduction in the quadratic graph term, not a change of asymptotic order. If  $B$  is allowed to scale with  $L$ , choosing  $B = \Theta(L^{1/3})$  balances the two terms and gives  $O(L^{4/3})$  edges.*

This edge count does not say the implementation is automatically faster. It says where the architecture moves communication: from dense token-token interactions to short token windows plus a smaller graph. For hidden size  $d$ , the leading attention arithmetic is  $O(MB^2d + M^2d) = O(LBd + L^2d/B^2)$ . The AST gate adds  $O(Ld_{\text{ast}}d)$  projection work and the MoE adds routed expert work. Those extra terms are exactly why the AST channel must later pass an efficiency ablation.

**Corollary 1** (Block-size trade-off). *If the dominant structured-attention cost is approximated by  $T(B) = aLB + bL^2/B^2$  for positive constants  $a$  and  $b$ , then the continuous minimizer is*

$$B^* = \left( \frac{2bL}{a} \right)^{1/3}.$$

*Thus no single fixed block size is theoretically privileged across all sequence lengths and hardware constants.*

This corollary is why I treat block size as an experimental variable rather than a magic number. The reference code supports arbitrary positive block sizes within shape constraints, and the paper should report a sweep before defending a particular  $B$ . The repository includes `scripts/architecture_cost_table.py`, which turns this accounting into a small JSON table. For example, at  $L = 2048, B = 64$ , dense causal attention has 2,098,176 causal token edges, while block-local causal attention plus the causal block graph has 67,088 counted edges in this simplified model. That number is not a runtime benchmark, but it does make the intended communication trade-off concrete.

**Proposition 4** (Long-range summary path). *For any two block indices  $i < j$ , one CSMT-GNN layer contains a differentiable computational path from the pooled state  $z_i$  to every injected token state  $H'_{j,b}$  in block  $j$ . A block-local Transformer layer without cross-block communication has no such one-layer path.*

This is the structural reason I added the graph. It does not force the model to use a distant block, but it makes the route available without asking token-level attention to pay the full  $L^2$  price.

**Proposition 5** (Residual AST gate stability). *Let  $g_{m,b} = 1 + \lambda \tanh(W_{\text{ast}} e_{m,b})$  with elementwise multiplication  $\bar{H}_{m,b} = H_{m,b} \odot g_{m,b}$ . If  $0 \leq \lambda < 1$ , then each coordinate of the gate lies in  $[1 - \lambda, 1 + \lambda]$ , and*

$$(1 - \lambda)\|H_{m,b}\|_2 \leq \|\bar{H}_{m,b}\|_2 \leq (1 + \lambda)\|H_{m,b}\|_2.$$

This is why I use a residual gate instead of a sigmoid gate. A sigmoid starts near  $1/2$  and immediately rescales the hidden stream. The residual gate starts near identity and lets the model learn how much syntax should matter.

**Proposition 6** (AST gate containment and cost boundary). *Let  $\mathcal{F}_0$  be the model family obtained by removing AST-gated pooling, and let  $\mathcal{F}_{\text{ast}}$  be the same family with the residual gate from Proposition 5. Because  $W_{\text{ast}} = 0$  gives  $g_{m,b} = 1$ ,  $\mathcal{F}_0 \subseteq \mathcal{F}_{\text{ast}}$ . Therefore the minimum unregularized empirical risk over  $\mathcal{F}_{\text{ast}}$  is no larger than the minimum over  $\mathcal{F}_0$ . If deployment is judged by  $J(f) = R(f) + \eta C(f)$ , where  $R$  is task risk,  $C$  is compute cost, and  $\eta > 0$ , then the AST side channel improves the deployment objective only when*

$$R(f_0^*) - R(f_{\text{ast}}^*) > \eta(C(f_{\text{ast}}^*) - C(f_0^*)).$$

This proposition is deliberately plain. It does not say optimization will find the better member of the larger family, and it does not say the added structure will generalize. It says the AST channel is not justified by expressivity alone. The extra projection work and preprocessing complexity must buy enough risk reduction to cross the displayed threshold.

**Proposition 7** (Prefix AST degradation metrics). *For a finished file, let  $A_t^{\text{full}}$  be the AST type id assigned to token  $t$  by a complete-file parse, and let  $A_t^{\text{pref}}$  be the id assigned when parsing only the generation prefix ending at  $t$ . Let  $F_t \in \{0, 1\}$  indicate that the prefix extractor fell back to a lexical token type, and let  $U_t \in \{0, 1\}$  indicate that the frozen training vocabulary maps the extracted type to `<UNKNOWN>`. Three quantities summarize the inference AST path:*

$$r_{\text{fb}} = \frac{1}{L} \sum_{t=1}^L F_t, \quad r_{\text{unk}} = \frac{1}{L} \sum_{t=1}^L U_t, \quad d_{\text{pref}} = \frac{1}{L} \sum_{t=1}^L \mathbf{1}[A_t^{\text{pref}} \neq A_t^{\text{full}}].$$

*If any of these rates is large, the AST side channel at inference time is not the same side channel used in offline training and must be cached, throttled, regularized against, or removed.*

**Proposition 8** (CVD as stochastic smoothing). *For an example with eligible definition blocks  $\mathcal{D}$ , let  $\xi_m \sim \text{Bernoulli}(p_{\text{cvd}})$  independently for  $m \in \mathcal{D}$ . Let  $f_\theta(x, A, \xi)$  denote the CSMT-GNN prediction after replacing value messages of sampled blocks by  $v_{\text{cvd}}$ . Training with CVD minimizes an unbiased Monte Carlo estimate of*

$$\mathcal{L}_{\text{cvd}}(\theta) = \mathbb{E}_{(x,y)} \mathbb{E}_\xi [\ell(f_\theta(x, A, \xi), y)].$$

*Moreover, a given eligible block is masked at least once in  $T$  independent presentations with probability  $1 - (1 - p_{\text{cvd}})^T$ .*



Table 2: Staged falsification plan. The point is not to make a small run look like a benchmark, but to decide what should be tested before scaling the architecture.

| Stage | Run                                                                                      | Pass signal                                                                                                                   | Stop signal                                                                                                              |
|-------|------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------|
| 0     | Static checks, prefix AST latency, fallback/unknown rates, shape tests.                  | No future-leakage path is detected; AST feature construction is measurable.                                                   | Prefix parsing dominates decoding or produces mostly fallback/unknown ids.                                               |
| 1     | Tiny diagnostic training on shadowing, imports, guarded attributes, and class ownership. | AST/CVD variants improve dependency-sensitive loss or preservation over token-only, random-dropout, and Transformer controls. | The full model fails to beat the cheap controls on diagnostics, or the parameter-neighbor Transformer explains the gain. |
| 2     | Small parameter-matched ablations on Python code.                                        | AST graph gains survive throughput and memory accounting.                                                                     | Gains are smaller than the cost of preprocessing and graph computation.                                                  |
| 3     | Larger code model attempt with fused kernels or grouped GEMM.                            | Structural gains persist when the reference bottlenecks are replaced by production kernels.                                   | The architecture only works in the toy setting or one hand-picked block size.                                            |

**Corollary 2** (Gradient exposure of the CVD replacement vector). *The replacement vector  $v_{\text{cvd}}$  receives gradient support only through positions whose value messages are sampled for replacement. If an example has  $D$  eligible blocks, at least one such position is sampled with probability  $1 - (1 - p_{\text{cvd}})^D$  on a single presentation and probability  $1 - (1 - p_{\text{cvd}})^{DT}$  across  $T$  independent presentations.*

This proposition is the right mathematical description of CVD. It is a structured stochastic smoothing objective over definition-bearing block messages. It is not a do-calculus intervention, and it gives no guarantee that structural hallucination will fall. What it does guarantee is that the training objective explicitly includes predictions under definition-message replacement, with a tunable exposure probability. It also explains why the old dummy-gradient trick was the wrong abstraction: the right statement is probabilistic gradient support, not guaranteed learning signal on every step. The reference trainer therefore leaves the optional L2 penalty on  $v_{\text{cvd}}$  disabled by default; turning it on is a practical stabilization choice, not part of the CVD objective above.

**Proposition 9** (Bounded CVD value perturbation). *Consider one graph-attention head at query block  $q$ , with attention weights  $a_{qm} \geq 0$  and  $\sum_m a_{qm} = 1$ . Let  $\mathcal{S}$  be the sampled CVD mask. If unmasked values and the replacement vector satisfy  $\|V_m\|_2 \leq R$  and  $\|v_{\text{cvd}}\|_2 \leq R$ , then replacing sampled values changes the head output by at most*

$$2R \sum_{m \in \mathcal{S}} a_{qm}.$$

This bound is another reason I prefer the CVD framing. The operation is a controlled value-message perturbation whose size depends on attention mass over masked definition blocks; it is not a literal deletion of a semantic variable.

## 6 Evaluation Protocol

I keep the evaluation plan separate from measured results. The theory above is a filter for internal consistency; it is not a substitute for empirical validation. Before treating CSMT-GNN as a competitive code model, I would start with five uncomfortable questions that can break the mechanism quickly.

Table 3: Recommended ablation matrix. The same names are used by the tiny diagnostic script before any larger parameter-matched run.

| Variant                             | AST gate | Block graph | CVD               | FFN path              |
|-------------------------------------|----------|-------------|-------------------|-----------------------|
| <code>transformer_baseline</code>   | No       | No          | No                | Transformer FFN       |
| <code>transformer_matched</code>    | No       | No          | No                | Wider Transformer FFN |
| <code>token_baseline</code>         | No       | No          | No                | Dense                 |
| <code>boundary_only</code>          | No       | No          | No                | Dense                 |
| <code>ast_only</code>               | Yes      | No          | No                | Dense                 |
| <code>graph_only</code>             | No       | Yes         | No                | Dense                 |
| <code>ast_graph</code>              | Yes      | Yes         | No                | Dense                 |
| <code>random_dropout_control</code> | Yes      | Yes         | Random blocks     | Dense                 |
| <code>variable_cvd</code>           | Yes      | Yes         | Definition blocks | Dense                 |
| <code>full_moe</code>               | Yes      | Yes         | Definition blocks | MoE                   |

**Can prefix AST features survive generation?** The first inference test measures whether the prefix AST path is usable during autoregressive decoding. Feed incomplete prefixes such as `def f(x:`, half-written class bodies, and open comprehensions into the prefix feature builder. The repository exposes this as a direct timing check through `inference_ast.py -repeat N` and as a full-vs-prefix degradation check through `scripts/prefix_ast_degradation.py`. Report parse success, fallback rate, unknown-token rate under the frozen training vocabulary, feature latency per generated token, and the divergence between prefix AST ids and complete-file AST ids after the code is finished. If this latency dominates decoding, AST conditioning should be cached, throttled to block boundaries, or removed.

**Does AST help enough to pay for itself?** The second ablation is deliberately harsh: train parameter-matched models with and without AST-gated pooling, and include a plain Transformer with comparable parameter count and training budget. Report tokens/sec, memory use, validation loss, HumanEval/MBPP pass@1, and AST-sensitive diagnostics. The decision rule should include efficiency: a small accuracy gain with large throughput loss is not a win for practical code modeling.

**How sensitive is the model to block size?** The third ablation sweeps  $B \in \{32, 64, 128\}$  and records throughput, memory, validation loss, and errors on examples where definitions or syntactic spans cross block boundaries. Compare boundary mixing on and off. A method that only works for one hand-picked block size is not robust enough to defend.

**Does CVD reduce structural hallucination?** The fourth ablation asks whether CVD is doing anything more useful than ordinary noise. Construct tests where the correct answer depends on definitions separated from uses. Examples include variable shadowing, renamed parameters, branch-guarded attributes, long-range imports, and class-method ownership. Compare no CVD, random block dropout, and variable-definition CVD. The key metric is not just surface match but whether generated code preserves the relevant dependency.

**Does the block graph carry useful state?** The fifth ablation is about the graph itself. Ablate graph attention, prefix injection, and graph depth. Inspect whether failures cluster in long files and whether the graph helps more as definition and use distance increases. This is where the architecture should earn its extra complexity.

**Minimum proof-of-concept before scale.** Before any large training run, I should run a deliberately small diagnostic: create synthetic Python files for variable shadowing, long-range imports, and guarded attributes, train tiny parameter-matched models for a fixed number of steps, and report whether AST/CVD changes the masked-token loss or dependency preservation on those diagnostics. The repository includes a diagnostic data generator for this purpose. Passing this test would not prove the architecture; failing it would be a useful early falsification.

## 7 Limitations

The first limitation is the validation level. The current artifact is a corrected prototype and paper design, not a fully trained model with large-scale empirical claims. I therefore separate mechanism checks from performance claims: prefix dependence, edge-count accounting, gate stability, and the stochastic objective induced by CVD help make the architecture inspectable, but they do not prove that CSMT-GNN beats established code LLMs. The release should be read as a falsifiable design artifact rather than a leaderboard result.

The second limitation is the causal framing. CVD was motivated by counterfactual thinking, but it is not an identified SCM intervention. If I or someone else wants to make a causal claim later, the work must specify variables, structural equations, intervention targets, and assumptions connecting code syntax to program behavior. I deliberately avoid that stronger claim here because the implementation does not earn it.

The third limitation is AST coverage. The current preprocessor focuses on Python. Extending to multilingual code requires language-specific parsers, stable vocabularies, and careful token alignment with the tokenizer used for training. Tree-sitter can help, but the engineering is not automatic.

The fourth limitation is inference-time parsing. Training can use offline AST features, while generation only has a partial prefix that may be syntactically invalid. I added a prefix AST builder with parser recovery and lexical fallback, but this is not a proof that the latency is acceptable. A real deployment must measure how often prefix features differ from complete-file features and whether caching at block boundaries is enough.

Finally, the architecture may simply be too expensive or too sensitive to the block size. I added an initially-zero boundary mixing module to reduce hard cuts through syntax, but the right block size still has to be measured. The design should survive cost-sensitive ablation before it is treated as a deployment candidate. If the AST gate and block graph do not improve downstream reliability enough to justify their cost, the right conclusion is to simplify the model.

## 8 Broader Impact

Better code models can make programming more accessible and can reduce routine engineering labor. They can also generate insecure code faster, help automate malicious scripting, or give users unwarranted confidence in plausible but wrong programs. CSMT-GNN is aimed at reducing structural mistakes in code generation, but the same capability could improve both benign and harmful automation. Any release of trained checkpoints should include standard code-safety evaluation, license-aware data handling, and clear warnings that generated code requires review and testing.

## 9 Conclusion

CSMT-GNN is my attempt to make a code language model carry a little more of the program’s structure without giving up token-level generation. The revised implementation is less grand than my first notes, and I think that is a good thing. It keeps the parts I still believe are worth testing—tokenizer-aligned AST ids, trainable syntax embeddings, prefix-masked block graph attention, and CVD over definition-bearing blocks—while removing claims and code paths that did not survive inspection. The result is narrower, but it is also more durable: a structure-aware code-model design, a reference implementation, and a diagnostic harness whose assumptions are visible enough to inspect and challenge.

## References

- Miltiadis Allamanis, Marc Brockschmidt, and Mahmoud Khademi. Learning to represent programs with graphs. In *International Conference on Learning Representations*, 2018.
- Uri Alon, Meital Zilberstein, Omer Levy, and Eran Yahav. code2vec: Learning distributed representations of code. In *Proceedings of the ACM on Programming Languages*, 2019.

- Marc Brockschmidt, Miltiadis Allamanis, Alexander L. Gaunt, and Oleksandr Polozov. Generative code modeling with graphs. In *International Conference on Learning Representations*, 2019.
- Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, Alex Ray, Raul Puri, Gretchen Krueger, Michael Petrov, Heidy Khlaaf, Girish Sastry, Pamela Mishkin, Brooke Chan, Scott Gray, Nick Ryder, Mikhail Pavlov, Alethea Power, Lukasz Kaiser, Mohammad Bavarian, Clemens Winter, Philippe Tillet, Felipe Petroski Such, Dave Cummings, Matthias Plappert, Fotios Chantzis, Elizabeth Barnes, Ariel Herbert-Voss, William Hebgen Guss, Alex Nichol, Alex Paino, Nikolas Tezak, Jie Tang, Igor Babuschkin, Suchir Balaji, Shantanu Jain, William Saunders, Christopher Hesse, Andrew N. Carr, Jan Leike, Josh Achiam, Vedant Misra, Evan Morikawa, Alec Radford, Matthew Knight, Miles Brundage, Mira Murati, Katie Mayer, Peter Welinder, Bob McGrew, Dario Amodei, Sam McCandlish, Ilya Sutskever, and Wojciech Zaremba. Evaluating large language models trained on code. In *arXiv preprint arXiv:2107.03374*, 2021.
- Tri Dao, Daniel Y. Fu, Stefano Ermon, Atri Rudra, and Christopher Ré. Flashattention: Fast and memory-efficient exact attention with io-awareness. In *Advances in Neural Information Processing Systems*, 2022.
- William Fedus, Barret Zoph, and Noam Shazeer. Switch transformers: Scaling to trillion parameter models with simple and efficient sparsity. *Journal of Machine Learning Research*, 23(120):1–39, 2022.
- Dmitry Lepikhin, HyukJoong Lee, Yuanzhong Xu, Dehao Chen, Orhan Firat, Yanping Huang, Maxim Krikun, Noam Shazeer, and Zhifeng Chen. Gshard: Scaling giant models with conditional computation and automatic sharding. In *International Conference on Learning Representations*, 2021.
- Erik Nijkamp, Bo Pang, Hiroaki Hayashi, Lifu Tu, Huan Wang, Yingbo Zhou, Silvio Savarese, and Caiming Xiong. Codegen: An open large language model for code with multi-turn program synthesis. *International Conference on Learning Representations*, 2023.
- Baptiste Roziere, Jonas Gehring, Fabian Gloeckle, Sten Sootla, Itai Gat, Xiaoqing Ellen Tan, Yossi Adi, Jingyu Liu, Tal Remez, Jérémy Rapin, Artyom Kozhevnikov, Ivan Evtimov, Joanna Bitton, Manish Bhatt, Cristian Canton Ferrer, Aaron Grattafiori, Wenhan Xiong, Alexandre Defossez, Jade Copet, Faisal Azhar, Hugo Touvron, Louis Martin, Nicolas Usunier, Thomas Scialom, and Gabriel Synnaeve. Code llama: Open foundation models for code. *arXiv preprint arXiv:2308.12950*, 2023.
- Noam Shazeer, Azalia Mirhoseini, Krzysztof Maziarczyk, Andy Davis, Quoc Le, Geoffrey Hinton, and Jeff Dean. Outrageously large neural networks: The sparsely-gated mixture-of-experts layer. In *International Conference on Learning Representations*, 2017.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems*, 2017.

## A Proofs for the Theoretical Checks

*Proof of Proposition 1.* Inside a block, causal self-attention prevents position  $b$  from reading positions greater than  $b$ . The block graph receives pooled summaries from valid blocks with index at most the current block; with the injection shift used in Section 3.5, block  $m$  receives only graph state from valid blocks  $< m$ . Padded block outputs are zeroed and are not injected into real token states. Therefore no later block can affect the current token state. The remaining possible leakage source is preprocessing: if AST ids are built from the complete file, a token can receive node information created by later syntax. The per-token prefix mode removes that source by construction, because the parser input ends at the token boundary. Composing these three facts gives the prefix dependence claim.  $\square$

*Proof of Proposition 2.* The update to block  $m$  depends on two quantities: the current token state  $\tilde{H}_{m,b}$  and the previous-block tail  $\tilde{H}_{m-1,B}$ . By causal block attention,  $\tilde{H}_{m,b}$  depends only on tokens within block  $m$  up to position  $b$ , and by the premise  $\tilde{H}_{m-1,B}$  depends only on tokens up to the end of block  $m-1$ . Their sum

therefore depends only on the union of those prefixes, which is exactly the token prefix ending at  $(m-1)B+b$ . The module contains no path from block  $m$  to block  $m-1$ , because the projection is applied from the previous block tail to the current block head only.  $\square$

*Proof of Proposition 3.* The full causal mask is triangular over  $L$  positions. Each of the  $M$  local blocks contributes  $B(B+1)/2$  causal edges. The prefix block graph is triangular over  $M$  block nodes, contributing  $M(M+1)/2$  edges. Dividing the structured edge count by the full edge count gives the displayed ratio. Substituting  $M = L/B$  gives  $O(LB) + O(L^2/B^2)$  edges. The two terms balance when  $LB = L^2/B^2$ , i.e.  $B = \Theta(L^{1/3})$ , yielding  $O(L^{4/3})$  edges.  $\square$

*Proof of Corollary 1.* Differentiate  $T(B)$  with respect to  $B$ :  $T'(B) = aL - 2bL^2/B^3$ . Setting  $T'(B) = 0$  gives  $B^3 = 2bL/a$ . The second derivative  $6bL^2/B^4$  is positive, so this critical point is a minimum.  $\square$

*Proof of Proposition 4.* In the prefix block graph, state  $s_{j-1}$  is computed by attention over block states  $z_1, \dots, z_{j-1}$ , which includes  $z_i$ . The graph-to-token injection then maps  $s_{j-1}$  through  $W_c$ , gates it with a token-dependent gate, and adds the result to every token in block  $j$ . This gives the path  $z_i \rightarrow s_{j-1} \rightarrow c_j \rightarrow H'_{j,b}$ . In a layer whose attention is restricted to tokens inside each block, the output of block  $j$  is a function only of block  $j$ 's input tokens, so no one-layer path from block  $i < j$  exists.  $\square$

*Proof of Proposition 5.* For every coordinate,  $\tanh(\cdot) \in [-1, 1]$ , so  $g_i \in [1 - \lambda, 1 + \lambda]$ . Since  $g_i$  is nonnegative when  $\lambda < 1$ , the squared norm of the gated vector is bounded between  $(1 - \lambda)^2 \sum_i H_i^2$  and  $(1 + \lambda)^2 \sum_i H_i^2$ . Taking square roots proves the claim.  $\square$

*Proof of Proposition 6.* Setting  $W_{\text{ast}} = 0$  makes the residual gate exactly one in every coordinate, so every no-AST model is represented inside the AST-gated family. Taking an infimum of the same empirical loss over a superset cannot increase the optimum. The deployment inequality follows by expanding  $J(f_{\text{ast}}^*) < J(f_0^*)$  and moving the compute term to the right-hand side.  $\square$

*Proof of Proposition 7.* The three rates are empirical averages over token-level Bernoulli indicators. They are therefore in  $[0, 1]$  and measure distinct failure modes: parser recovery or lexical fallback, vocabulary mismatch, and disagreement with a complete-file parse after the file is known. A model trained on complete-file or low-fallback AST ids receives a shifted side input when these rates grow. The proposition is a measurement contract rather than a performance guarantee: it identifies when inference-time AST conditioning has drifted from the training artifact.  $\square$

*Proof of Proposition 8.* The implementation samples  $\xi$  and evaluates the loss for that sampled masked graph. The stochastic gradient is therefore the gradient of one draw from the inner expectation. Under the usual interchange conditions used for stochastic gradient training, its expectation equals  $\nabla_{\theta} \mathcal{L}_{\text{cvd}}(\theta)$ . The second statement is the complement of never sampling the block in  $T$  independent Bernoulli trials.  $\square$

*Proof of Corollary 2.* The vector  $v_{\text{cvd}}$  is used only when some Bernoulli mask is one. The probability that all  $D$  masks are zero is  $(1 - p_{\text{cvd}})^D$ , and the same complement argument over  $DT$  independent trials gives the multi-pass statement.  $\square$

*Proof of Proposition 9.* The unmasked head output is  $o_q = \sum_m a_{qm} V_m$  and the masked output is  $\tilde{o}_q = \sum_{m \notin \mathcal{S}} a_{qm} V_m + \sum_{m \in \mathcal{S}} a_{qm} v_{\text{cvd}}$ . Their difference is  $\tilde{o}_q - o_q = \sum_{m \in \mathcal{S}} a_{qm} (v_{\text{cvd}} - V_m)$ . By the triangle inequality and the norm bounds,  $\|\tilde{o}_q - o_q\|_2 \leq \sum_{m \in \mathcal{S}} a_{qm} (\|v_{\text{cvd}}\|_2 + \|V_m\|_2) \leq 2R \sum_{m \in \mathcal{S}} a_{qm}$ .  $\square$

## B Reproducibility Notes

The public reproducibility repository is available at <https://github.com/ShengyaoLiang/csmt-gnn>. The current v0.1.1 archived release is <https://doi.org/10.5281/zenodo.20844185>; the stable all-version Zenodo concept DOI is <https://doi.org/10.5281/zenodo.20840624>. I release the source package, paper source, tests, small diagnostic artifacts, and packaging scripts so that the evidence chain can be inspected without relying on private training data, hidden scripts, or unpublished checkpoints.

The source tree contains:

- `ast_preprocessor.py`: offline AST id and variable-mask builder.
- `inference_ast.py`: prefix AST feature builder for incomplete generation prefixes.
- `scripts/prefix_ast_degradation.py`: full-vs-prefix AST drift measurement.
- `scripts/architecture_cost_table.py`: structural edge-count and block-size trade-off audit.
- `diagnostics.py`: tiny structural diagnostic data generator.
- `scripts/structural_probe_eval.py`: definition-use coverage check for the diagnostic cases.
- `scripts/cvd_mask_audit.py`: CVD eligible/sampled block audit.
- `scripts/validate_data_pipeline.py`: NumPy token/AST preflight before training fast paths.
- `scripts/diagnostic_poc_train.py`: tiny ablation sanity check.
- `transformer_baseline.py`: independent token-only Transformer baseline for the diagnostic comparison.
- `scripts/build_arxiv_package.py`: arXiv source package builder.
- `csmt_gnn.py`: model architecture.
- `train.py`: single-process and distributed training entry point.
- `requirements.txt`: minimal Python dependencies.
- `results/lowcompute_validation_summary.md`: local laptop smoke validation record, if included in the public code release.

A minimal debugging command is:

```
python ast_preprocessor.py \
  --source-file example.py \
  --output-dir ast_data
```

For training artifacts, use tokenizer offsets. The default is complete-file offline parsing:

```
python ast_preprocessor.py \
  --source-file example.py \
  --output-dir ast_data \
  --tokenizer-name-or-path /path/to/tokenizer
```

For strict prefix-only preprocessing studies, use the explicit slow mode:

```
python ast_preprocessor.py \
  --source-file example.py \
  --output-dir ast_data_prefix \
  --tokenizer-name-or-path /path/to/tokenizer \
  --per-token-prefix
```

For prefix inference features:

```
python inference_ast.py --source-file partial_generation.py \
  --vocab-path ast_data/ast_vocab.json \
  --repeat 32
```

For full-vs-prefix AST degradation:

```
python scripts/prefix_ast_degradation.py \  
  --source-file example.py \  
  --vocab-path ast_data/ast_vocab.json \  
  --output results/prefix_ast_degradation.json
```

For the structural edge-count audit:

```
python scripts/architecture_cost_table.py \  
  --output results/architecture_cost_table.json
```

For the smallest structural diagnostic arrays:

```
python diagnostics.py --output-dir tmp/diagnostics --block-size 8
```

For structural coverage and CVD sampling audits:

```
python scripts/structural_probe_eval.py \  
  --output results/structural_probe_eval.json \  
  --block-size 8 --max-tokens 64
```

```
python scripts/cvd_mask_audit.py \  
  --output results/cvd_mask_audit.json \  
  --steps 24 --hidden-size 16 --ast-dim 8 \  
  --block-size 8 --max-tokens 64
```

A minimal training command, assuming paired token arrays are present, is:

```
python train.py \  
  --data-path tokenized_data \  
  --ast-path ast_data \  
  --num-layers 12 \  
  --hidden-size 768 \  
  --block-size 64 \  
  --max-tokens 2048 \  
  --boundary-width 2
```

When PyTorch is available, the local diagnostic sanity check is:

```
python scripts/diagnostic_poc_train.py --steps 12 \  
  --output results/diagnostic_poc_transformer.json
```

On my local laptop, I ran this diagnostic at `hidden_size=16`, `max_tokens=64`, and 12 steps per variant, plus small block-size and seed smoke sweeps. The diagnostic now includes a pure token-only Transformer and a rough parameter-neighbor Transformer control. The unit tests now pass at 39/39 after adding fail-fast checks for the Transformer controls, incremental AST configuration checks, opt-in CVD audit coverage, validated-pipeline input range controls, graph padding-mask coverage, trainer prefix-trim coverage, diagnostic prefix-audit coverage, and a regression test for token-level CVD masks on short next-token inputs. A two-step diagnostic run also recorded `prefix_feature_audit.all_prefix_aligned=1.0` with an empty violation list, confirming that the diagnostic CSMT variants receive prefix-clipped AST ids and definition masks. The data-pipeline preflight also passed on the tiny diagnostic arrays, checking token ids, AST ids, block shapes, masks, and usable sequence lengths. The structural probe found cross-block definition-use pairs in all three diagnostic cases. The CVD audit recorded 72 variable-eligible blocks and 112 random-control eligible blocks over 24 audited steps, with both sampling rates near the configured 0.2. The tiny training run showed that random CVD and variable CVD were almost identical in final next-token loss. I treat this record only as a mechanism and plumbing check: it shows that the reference paths can train on CPU and that the masks are observable, not that the architecture is better than a baseline.

My reading of these numbers is cautious. On the seed-7 run, block-graph variants improve the cross-block use loss over the CSMT token baseline. Across seeds 1–3, however, the rough parameter-neighbor Transformer

Table 4: Local CPU diagnostic at seed 7, block size 8, hidden size 16, and 12 steps per variant. This is not a benchmark; it is a small falsification and plumbing check.

| Variant                     | Params | Final loss | Eval loss | Cross-block use loss |
|-----------------------------|--------|------------|-----------|----------------------|
| <b>transformer_baseline</b> | 4688   | 3.6034     | 3.6342    | 3.6809               |
| <b>transformer_matched</b>  | 6992   | 3.6751     | 3.5889    | 3.7469               |
| <b>token_baseline</b>       | 4576   | 3.7727     | 3.7131    | 3.9152               |
| <b>graph_only</b>           | 6384   | 3.7367     | 3.6630    | 3.6176               |
| <b>ast_graph</b>            | 6960   | 3.9444     | 3.6480    | 3.5019               |
| <b>variable_cvd</b>         | 6976   | 3.6431     | 3.5470    | 3.8563               |
| <b>full_moe</b>             | 8160   | 3.6942     | 3.6087    | 3.2707               |

Table 5: Seed 1–3 mean local diagnostic results at block size 8. Parentheses show population standard deviation. The close result between **transformer\_matched** and **ast\_graph** is why I do not claim a Transformer win or a CSMT-GNN win from this run.

| Variant                     | Params | Final loss      | Cross-block use loss |
|-----------------------------|--------|-----------------|----------------------|
| <b>transformer_baseline</b> | 4688   | 3.7680 (0.0121) | 3.6739 (0.1213)      |
| <b>transformer_matched</b>  | 6992   | 3.6885 (0.0832) | 3.5243 (0.3131)      |
| <b>ast_graph</b>            | 6960   | 3.6777 (0.0955) | 3.5366 (0.2761)      |
| <b>variable_cvd</b>         | 6976   | 3.7446 (0.0514) | 3.8471 (0.2794)      |
| <b>full_moe</b>             | 8160   | 3.6761 (0.0412) | 3.7173 (0.1219)      |

and **ast\_graph** are essentially tied on mean cross-block use loss. CVD is also unstable in this tiny setup. This is exactly the kind of result I wanted the diagnostic to be able to show: the architecture is now comparable against a real Transformer control, but it has not earned a claim that it beats Transformer.

The arXiv source package is built with:

```
python scripts/build_arxiv_package.py
```